

How does the GAN make $p_{\text{data}} \approx p_{\text{model}}$?

The Generator Loss function we constructed is a proxy to achieve the goal

$$p_{\text{model}} \approx p_{\text{data}}$$

That is: the distribution of samples produced by the Generator is (approximately) the same as the "true" distribution

- we note that we don't know the "true" p_{data}
 - we only have available a sample and those the training set defines an *empirical* distribution

There are several ways to quantify

$$p_{\text{model}} \approx p_{\text{data}}$$

One choice would be the minimization of KL Divergence

- $\mathbf{KL}(p_{\text{data}} || p_{\text{model}})$

As a reminder: we now show that this is equivalent to Maximum Likelihood estimation

Choose p_{model} to Minimize

$$\begin{aligned}\mathbf{KL}(p_{\text{data}} || p_{\text{model}}) &= \int_{\mathbf{x}} p_{\text{data}}(\mathbf{x}) \left(\log \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{model}}(\mathbf{x})} \right) d\mathbf{x} && \text{Definition of KL} \\ &= \mathbb{E}_{\mathbf{x} \in p_{\text{data}}} \log(p_{\text{data}}(\mathbf{x})) - \log(p_{\text{model}}(\mathbf{x})) && \text{Definition of log} \\ \text{minimizing KL} &&& \\ &\approx \mathbb{E}_{\mathbf{x} \in p_{\text{data}}} - \log(p_{\text{model}}(\mathbf{x})) && \text{Since } \log(p_{\text{data}}(\mathbf{x})) \text{ is constant}\end{aligned}$$

So minimizing \mathbf{KL} is equivalent to minimizing the Negative Log Likelihood.

Notice that the expectation is over the "true" distribution p_{data} .

The expectation is certainly reasonable for training but perhaps not best for the purposes of generating synthetic data

- Measures fidelity to training data
- NOT how "realistic" the synthetic data is
- the penalty for p_{model} placing large probability mass around a particular $\hat{\mathbf{x}}'$ is small when $p_{\text{data}}(\hat{\mathbf{x}}') \approx 0$
 - so Generator may create large quantity of synthetic data that is improbable given the training set

If we knew the true p_{data} , a better objective to minimize for the purpose of generating synthetic data would be the similar

$$\mathbf{KL}(p_{\text{model}} || p_{\text{data}})$$

which is equivalent to maximizing

$$\mathbb{E}_{\mathbf{x} \in p_{\text{model}}} - \log(p_{\text{data}}(\mathbf{x}))$$

The expectation is over the synthetic data, not the true data

- $\log(p_{\text{data}}(\mathbf{x}))$ is defined as log of Perplexity
 - an element of "surprise" in seeing \mathbf{x}
- So the expectation asks: for each synthetic datum generated, how likely is it to occur in the true distribution ?

This is merely a theoretical argument

- In practical terms: we only have empirical p_{data}

So can't evaluate

- log Perplexity

$$p_{\text{data}}(\hat{\mathbf{x}})$$

for $\hat{\mathbf{x}} \in p_{\text{model}}$

- unless synthetic $\hat{\mathbf{x}}$ replicates a sample in the training data

Jensen-Shannon Divergence

We have observed that the KL divergence is *not* symmetric

$$\mathbf{KL}(P||Q) \neq \mathbf{KL}(Q||P)$$

because the expectations are taken over different distributions.

An alternative measure of similarity of two distributions is the Jensen-Shannon Divergence (JSD)

$$\begin{aligned}\text{JSD}(P||Q) &= \text{JSD}(Q||P) \\ &= \frac{1}{2} \text{KL} \left(P || \frac{P+Q}{2} \right) + \\ &\quad \frac{1}{2} \text{KL} \left(Q || \frac{P+Q}{2} \right)\end{aligned}$$

This measure is

- symmetric
- is a kind of mixture of $\mathbf{KL}(P||Q)$ and $\mathbf{KL}(Q||P)$.

Huszar (<https://arxiv.org/pdf/1511.05101.pdf>) has a Generalized JSD which interpolates between the two terms

$$\begin{aligned}\text{JSD}_{\pi}(P||Q) &= \text{JSD}(Q||P) \\ &= \pi \text{KL}(P || \pi P + (1 - \pi)Q) + \\ &\quad (1 - \pi) \text{KL}(Q || \pi P + (1 - \pi)Q)\end{aligned}$$

The Generalized JSD

- **Not** symmetric although
$$\text{JSD}_{\pi}(P||Q) = \text{JSD}_{1-\pi}(Q||P)$$

Huszar shows that, for small values of π

$$\frac{\text{JSD}_{\pi}(P||Q)}{\pi} \approx \text{KL} (P || Q)$$

and

$$\frac{\text{JSD}_{1-\pi}(P||Q)}{1 - \pi} \approx \text{KL} (Q || P)$$

In the first case

- $\text{JSD}_{\pi}(P||Q)$ is proportional to Maximum Likelihood

In the second case

- $\text{JSD}_{1-\pi}(P||Q)$ is proportional to $\text{KL} (Q || P)$

In implementing Generalized JSD

- The Discriminator is trained (as usual) on a mix of real and fake examples
 - But *not* in equal numbers
 - π is fraction of samples from Q
 - $(1 - \pi)$ is fraction of samples from P
 - $\pi < \frac{1}{2}$: real samples over represented
 - $\pi > \frac{1}{2}$: biased toward Q
- Explains why we often see training with Generator updated twice for each update of Discriminator ?

Adversarial Training and the Jensen-Shannon Divergence

The Discriminator Loss \mathcal{L}_D

- summed over all examples
 - (ignoring the $\frac{1}{2}$ from the previous presentation where we assumed equal number of Real and Fake)

is

$$\mathcal{L}_D = - \left(\mathbb{E}_{\mathbf{x}^{(i)} \in p_{\text{data}}} \log D(\mathbf{x}^{(i)}) + \mathbb{E}_{\mathbf{x}^{(i)} \in p_{\text{model}}} \log(1 - D(\mathbf{x}^{(i)})) \right) \quad D(G(\mathbf{z})) = \mathbf{x}^f$$

We also showed that the optimal Discriminator results in

$$D^*(\mathbf{x}) = \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{model}}(\mathbf{x}) + p_{\text{data}}(\mathbf{x})}$$

Plugging $D^*(\mathbf{x})$ into \mathcal{L}_D (Goodfellow Equation):

$$\begin{aligned}
 \mathcal{L}_D &= - \left(\mathbb{E}_{\mathbf{x}^{(i)} \in p_{\text{data}}} \log \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{model}}(\mathbf{x}) + p_{\text{data}}(\mathbf{x})} + \mathbb{E}_{\mathbf{x}^{(i)} \in p_{\text{model}}} \log \frac{p_{\text{model}}(\mathbf{x})}{p_{\text{model}}(\mathbf{x}) + p_{\text{data}}(\mathbf{x})} \right) \\
 &= - (\mathbf{KL}(p_{\text{data}} || p_{\text{model}}(\mathbf{x}) + p_{\text{data}}(\mathbf{x})) + \mathbf{KL}(p_{\text{model}} || p_{\text{model}}(\mathbf{x}) + p_{\text{data}}(\mathbf{x}))) \\
 &= - \left(\log 4 + \mathbf{KL}(p_{\text{data}} || \frac{p_{\text{model}}(\mathbf{x}) + p_{\text{data}}(\mathbf{x})}{2}) + \mathbf{KL}(p_{\text{model}} || \frac{p_{\text{model}}(\mathbf{x}) + p_{\text{data}}(\mathbf{x})}{2}) \right) \\
 \\
 &= - (\log 4 + 2 * \text{JSD}(p_{\text{data}} || p_{\text{model}}))
 \end{aligned}$$

The above equations shows that

- minimizing KL Divergence (second line above)
- under the assumption that the Discriminator can train to be the **optimal** adversary

results in \mathcal{L}_D becoming equivalent to Jensen-Shannon Distance (last line above)

So solving the minimax optimally minimizes the JSD divergence between p_{data} and p_{model} .

In [2]: `print("Done")`

Done

