

Hopefully by now, we know about *Gradient Descent*

- Solving for weights/parameters
- That minimize a loss function
- By updating weights/parameters in the *negative* direction of the gradients with respect to the parameters/weights

In code, it looks like this

- from [Keras docs \(https://colab.research.google.com/github/keras-team/keras-io/blob/master/guides/ipynb/customizing_what_happens_in_fit.ipynb#scrollTo=9z4\)](https://colab.research.google.com/github/keras-team/keras-io/blob/master/guides/ipynb/customizing_what_happens_in_fit.ipynb#scrollTo=9z4)
- one step of Gradient Descent (inputs are a mini-batch of examples)

```
with tf.GradientTape() as tape:
    y_pred = self(x, training=True) # Forward pass
    # Compute the loss value
    # (the loss function is configured in `compile()`)
    loss = self.compiled_loss(y, y_pred, regularization_losses=self.losses)

# Compute gradients
trainable_vars = self.trainable_variables
gradients = tape.gradient(loss, trainable_vars)

# Update weights
self.optimizer.apply_gradients(zip(gradients, trainable_vars))
```

Key points

- Define a loss \mathcal{L}
 - the loss is dependent on the weights ("trainable variables") of the model
- Compute the loss within the scope of `tf.GradientTape()`
 - Enables TensorFlow to compute gradients of any variable accessed in the scope
 - Loss calculated via `self.compiled_loss` in this case
 - but any calculation that you would chose to define
- Obtain the gradients of the loss with respect to the trainable variables
- Updates the trainable variables
 - `self.optimizer.apply_gradients(zip(gradients, trainable_vars))` in this case
 - General case `weight += - learning_rate * gradient`
 - Subtract the gradient: we are descending (reducing loss)

Gradient Ascent is nearly identical

- Except that we update
 - a collection of variables
 - **not necessarily** the weights
 - perhaps some other variable
 - in the *positive* direction of the gradients
- So as to *maximize* a function ("utility")
 - we will continue, in code, to use "loss" for the function/variable name

In code, it looks like this:

```
with tf.GradientTape() as tape:
    tape.watch(vars)
    loss = compute_loss(vars)

# Compute gradients.
gradients = tape.gradient(loss, vars)

vars += learning_rate * gradients
```

- `vars` is a list of variables
- loss is dependent on `vars`
- we compute the gradient of the loss with respect to `vars`
- we *add* the gradient wrt `vars`:
 - we are ascending (increasing loss: better to call it "utility")

Uses of Gradient Ascent

We will show some interesting things you can do using Gradient Ascent.

We need to identify

- the property being maximized
- the variables that will be adjusted by the maximizer

Our examples identify

- the property as defined by variable \mathcal{L} (which we are maximizing, not minimizing)
- the variables being adjusted are the **inputs** to the NN
 - $\text{vars} = \mathbf{y}_{(0)}$

That is

- we are solving for the inputs that maximize a property.

Property to maximize: value of a single "logit" of the Classifier head

Suppose our Neural Network \mathbb{C} terminates in a Classifier head, over classes $\{c_1, \dots, c_k\}$.

The Classifier Head is a Dense layer with k units ("logits"), one per class.

For example, MNIST digit classification

- input is a (28×28) grid of pixel values
- there are 10 "logits"
 - corresponding to each of the digits $\{0, 1, \dots, 9\}$

Define the property to be maximized

- The value of logit corresponding to c_j

Gradient Ascent will find the input value to \mathbb{C} that will be classified with highest probability as being from class c_j .

This is the "paradigmatic" input of class c_j .

Using MNIST digit classification as our example.

Suppose we choose the property: "Maximize the logit for digit 8"

- we want to solve for a (28×28) pixel grid
- that is classified as an "8" with highest probability

Note

- the pixel grid solution *does not necessarily* look like a digit !

Property to maximize: summary of values of one feature map

We can generalize the MNIST example

- suppose we have a multi-layer Sequential NN
- we want to *interpret* the purpose of layer l
 - in our first case, $l = L$ (the Classifier layer)
 - here: l can be an intermediate layer

Recall that a *feature map* is

- a Tensor (with shape equal to the spatial dimensions)
- corresponding to the value of a single feature at some layer l
 - over each spatial location

Since this feature is not a singleton, imagine we reduced it to a single value

- e.g., maximum value

Define the property to be maximized as

- a single value that summarizes *all* the values of the single feature map at layer l
 - for example: the maximum value over the Tensor

We use Gradient Ascent

- to solve for a (28×28) pixel grid
- the *maximally activates* the feature map at layer l

Again, the solution pixel grid may not look like a digit

- but it may reveal a pattern that "triggers" this layer
 - for example: "strong vertical line"
 - indicative of digits 1, 4, 9

This is one way in which we attempt to discern what role each layer serves.

Visualizing what convnets learn, via Gradient Ascent

Let's illustrate Gradient Ascent to visualize what one feature map within a Convolutional Layer of an Image Classifier is "looking for"

[Visualizing what convnets learn \(https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/vision/ipynb/visualizing_what_convnets_learn.ipynb#\)](https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/vision/ipynb/visualizing_what_convnets_learn.ipynb#)

A blog post from a [previous version] of the code shows the patterns of multiple feature maps at multiple layers.

In [2]: `print("Done")`

Done

