

Tokenizers

The first step in the NLP pipeline is converting sequences of text into sequences of tokens.

This process is called *Tokenization*.

Fortunately, many toolkits provide implementations of Tokenizers so that you don't have to do this yourself.

Our concern is to discuss the concepts and limitations.

We refer the interested read to some good *practical* resources on Tokenization

- [Hugging Face course section on Tokenization](https://huggingface.co/course/chapter2/4?fw=pt)
(<https://huggingface.co/course/chapter2/4?fw=pt>).
- [TensorFlow tensorflow_text API](https://www.tensorflow.org/text) (<https://www.tensorflow.org/text>).
- [Keras Project NLP API](https://keras.io/api/keras_nlp/tokenizers/) (https://keras.io/api/keras_nlp/tokenizers/).
 - Note: this is the Keras project **not** the implementation of Keras in TensorFlow

Word based tokenization

There are several varieties of Tokenization.

The first is Word Based

- convert words into tokens

It would be naive to imagine that each word is a token

- The size of the vocabulary **V** can easily be tens of thousands.
- Not only are there many words, there are *variations* of each word
 - verbs: past, present and future tense
 - nouns:
 - singular, plural
 - gender agreement with subject

"Classical" NLP has developed techniques to eliminate word variation by conversion to "canonical" form

- stemming: convert word to its root form (root may not be in dictionary)
 - drive, driver, driving \mapsto driv
- lemmatization: convert word to a root that is in dictionary
 - him: \mapsto he
 - took \mapsto take
 - earlier \mapsto early

So one way to reduce the vocabulary size is by using these techniques.

There is still a place for this techniques in Deep Learning.

- [spaCy](https://spacy.io/) (<https://spacy.io/>) is a very popular toolkit for dealing with text.

It is also possible that Embeddings will create a small number of "dimension" (e.g., variations) that are common to many words

- the "plural" dimension

Issues with Word based tokenization

Most word based tokenizers assumed that words are delimited by the space and a fixed number of punctuation characters.

A big unresolved issue with word based tokenization:

- No matter how large the vocabulary \mathbf{V}
- There will still be words that are Out of Vocabulary (OOV)

One way to address this is with Character based tokenization

- each character is a token

But this may inhibit learning

- Characters are too low level
- Compared to words

Sub-word Tokenization

An elegant solution to the OOV issue is *sub-word tokenization*

- break an OOV token into pieces, each of which is in-vocabulary
 - here's \mapsto here, \', s
 - tokenizer \mapsto token, ##izer

Some of the objectives are

- to **not** split frequently occurring words into pieces
 - So that the model consuming the tokens has access to meaningful single word pieces
- to recognize pieces that are common prefix and suffix forms
 - So that the model consuming the tokens can derive the commonality indicated by the suffix
 - e.g. "izer"

Sub-word tokenization has become fairly common in modern NLP systems

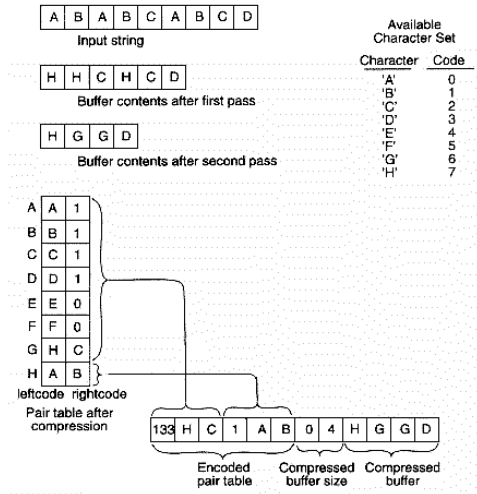
- Balances brevity of character encoding
- With expressiveness of word encoding

Sub-word tokenization: Byte Pair Encoding (BPE)

There are several methods to achieve sub-word Tokenization.

A simple version is derived from [Byte Pair Encoding \(BPE\)](https://arxiv.org/pdf/1508.07909.pdf).

- BPE is a way to replace common sequences of bytes with a short symbol
- Word segmentation: the analog of BPE, where we replace common sequences of tokens with a short symbol



Picture from: <http://twimags.com/ddj/cuj/images/cuj9402gage/fig1.gif>

- Pass 1: AB \mapsto H
 - byte sequence "A", "B" occurs 3 times in input
 - replaced with symbol "H" (recorded dictionary entry "H" _)
- Pass 2: HC \mapsto G
 - symbol sequence "H", "C" (occurring twice) replaced with symbol "G"

WordPiece sub-word tokenization

WordPiece (<https://arxiv.org/pdf/1609.08144.pdf>) is a variant of BPE.

Sentence Piece sub-word tokenization

BPE provides a unique encoding (*segmentation*) of a sentence of text.

But there are multiple ways of dividing words into segments.

Subwords (· means spaces)	Vocabulary id sequence
·Hell/o/·world	13586 137 255
·H/ello/·world	320 7363 255
·He/llo/·world	579 10115 255
·/He/l/o/·world	7 18085 356 356 137 255
·H/el/l/o/·world	320 585 356 137 7 12295

Table 1: Multiple subword sequences encoding the same sentence “Hello World”

Attribution: Table 1 in <https://arxiv.org/pdf/1804.10959.pdf>

The paper [Subword Regularization: Improving Neural Network Translation Models with Multiple Subword Candidates \(https://arxiv.org/pdf/1804.10959.pdf\)](https://arxiv.org/pdf/1804.10959.pdf) proposed

- using multiple segmentations
- associating a probability with each

The advantages include

- accommodation of multiple languages
- increased robustness of training

[SentencePiece \(https://arxiv.org/pdf/1808.06226.pdf\)](https://arxiv.org/pdf/1808.06226.pdf) is an Open Source implementation of the subword regularization technique.

```
In [1]: print("Done")
```

Done

